

```

1 import numpy as np # linear algebra
2 import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
3 from scipy import stats
4 import matplotlib.pyplot as plt
5 import seaborn as sns

```

```

1 csv_path = "yulu dataset.txt"
2 df = pd.read_csv(csv_path, delimiter=",")
3 df.head()

```

```

↗
      datetime season holiday workingday weather temp atemp humidity windspeed casual registered count
0 2011-01-01 00:00:00      1      0          0      1  9.84 14.395      81        0.0      3         13      16
1 2011-01-01 01:00:00      1      0          0      1  9.02 13.635      80        0.0      8         32      40
2 2011-01-01 02:00:00      1      0          0      1  9.02 13.635      80        0.0      5         27      32
3 2011-01-01 03:00:00      1      0          0      1  9.84 14.395      75        0.0      3         10      13
4 2011-01-01 04:00:00      1      0          0      1  9.84 14.395      75        0.0      0          1       1

```

```

1 # no of rows amd columns in dataset
2 print(f"# rows: {df.shape[0]} \n# columns: {df.shape[1]}")

```

```

↗ # rows: 10886
# columns: 12

```

```
1 df.info()
```

```

↗ <class 'pandas.core.frame.DataFrame'>
RangeIndex: 10886 entries, 0 to 10885
Data columns (total 12 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   datetime    10886 non-null  object
1   season      10886 non-null  int64
2   holiday     10886 non-null  int64
3   workingday  10886 non-null  int64
4   weather     10886 non-null  int64
5   temp        10886 non-null  float64
6   atemp       10886 non-null  float64
7   humidity    10886 non-null  int64
8   windspeed   10886 non-null  float64
9   casual      10886 non-null  int64
10  registered  10886 non-null  int64
11  count       10886 non-null  int64
dtypes: float64(3), int64(8), object(1)
memory usage: 1020.7+ KB

```

```

1 df['datetime'] = pd.to_datetime(df['datetime'])
2
3 cat_cols= ['season', 'holiday', 'workingday', 'weather']
4 for col in cat_cols:
5     df[col] = df[col].astype('object')

```

```
1 df.info()
```

```

↗ <class 'pandas.core.frame.DataFrame'>
RangeIndex: 10886 entries, 0 to 10885
Data columns (total 12 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   datetime    10886 non-null  datetime64[ns]
1   season      10886 non-null  object
2   holiday     10886 non-null  object
3   workingday  10886 non-null  object
4   weather     10886 non-null  object
5   temp        10886 non-null  float64
6   atemp       10886 non-null  float64
7   humidity    10886 non-null  int64
8   windspeed   10886 non-null  float64
9   casual      10886 non-null  int64
10  registered  10886 non-null  int64
11  count       10886 non-null  int64
dtypes: datetime64[ns](1), float64(3), int64(4), object(4)
memory usage: 1020.7+ KB

```

```
1 df.iloc[:, 1:].describe(include='all')
```

↗

	season	holiday	workingday	weather	temp	atemp	humidity	windspeed	casual	registered	
count	10886.0	10886.0	10886.0	10886.0	10886.00000	10886.000000	10886.000000	10886.000000	10886.000000	10886.000000	108
unique	4.0	2.0	2.0	4.0	NaN	NaN	NaN	NaN	NaN	NaN	
top	4.0	0.0	1.0	1.0	NaN	NaN	NaN	NaN	NaN	NaN	
freq	2734.0	10575.0	7412.0	7192.0	NaN	NaN	NaN	NaN	NaN	NaN	
mean	NaN	NaN	NaN	NaN	20.23086	23.655084	61.886460	12.799395	36.021955	155.552177	1
std	NaN	NaN	NaN	NaN	7.79159	8.474601	19.245033	8.164537	49.960477	151.039033	1
min	NaN	NaN	NaN	NaN	0.82000	0.760000	0.000000	0.000000	0.000000	0.000000	
25%	NaN	NaN	NaN	NaN	13.94000	16.665000	47.000000	7.001500	4.000000	36.000000	
50%	NaN	NaN	NaN	NaN	20.50000	24.240000	62.000000	12.998000	17.000000	118.000000	1
75%	NaN	NaN	NaN	NaN	26.24000	31.060000	77.000000	16.997900	49.000000	222.000000	2
max	NaN	NaN	NaN	NaN	41.00000	45.455000	100.000000	56.996900	367.000000	886.000000	9

```
1 # detecting missing values in the dataset
2 df.isnull().sum()
```

↗

datetime	0
season	0
holiday	0
workingday	0
weather	0
temp	0
atemp	0
humidity	0
windspeed	0
casual	0
registered	0
count	0
dtype: int64	

```
1 # minimum datetime and maximum datetime
2 print(df['datetime'].min(), df['datetime'].max())
3 # number of unique values in each categorical columns
4 df[cat_cols].melt().groupby(['variable', 'value'])[['value']].count()
```

↗

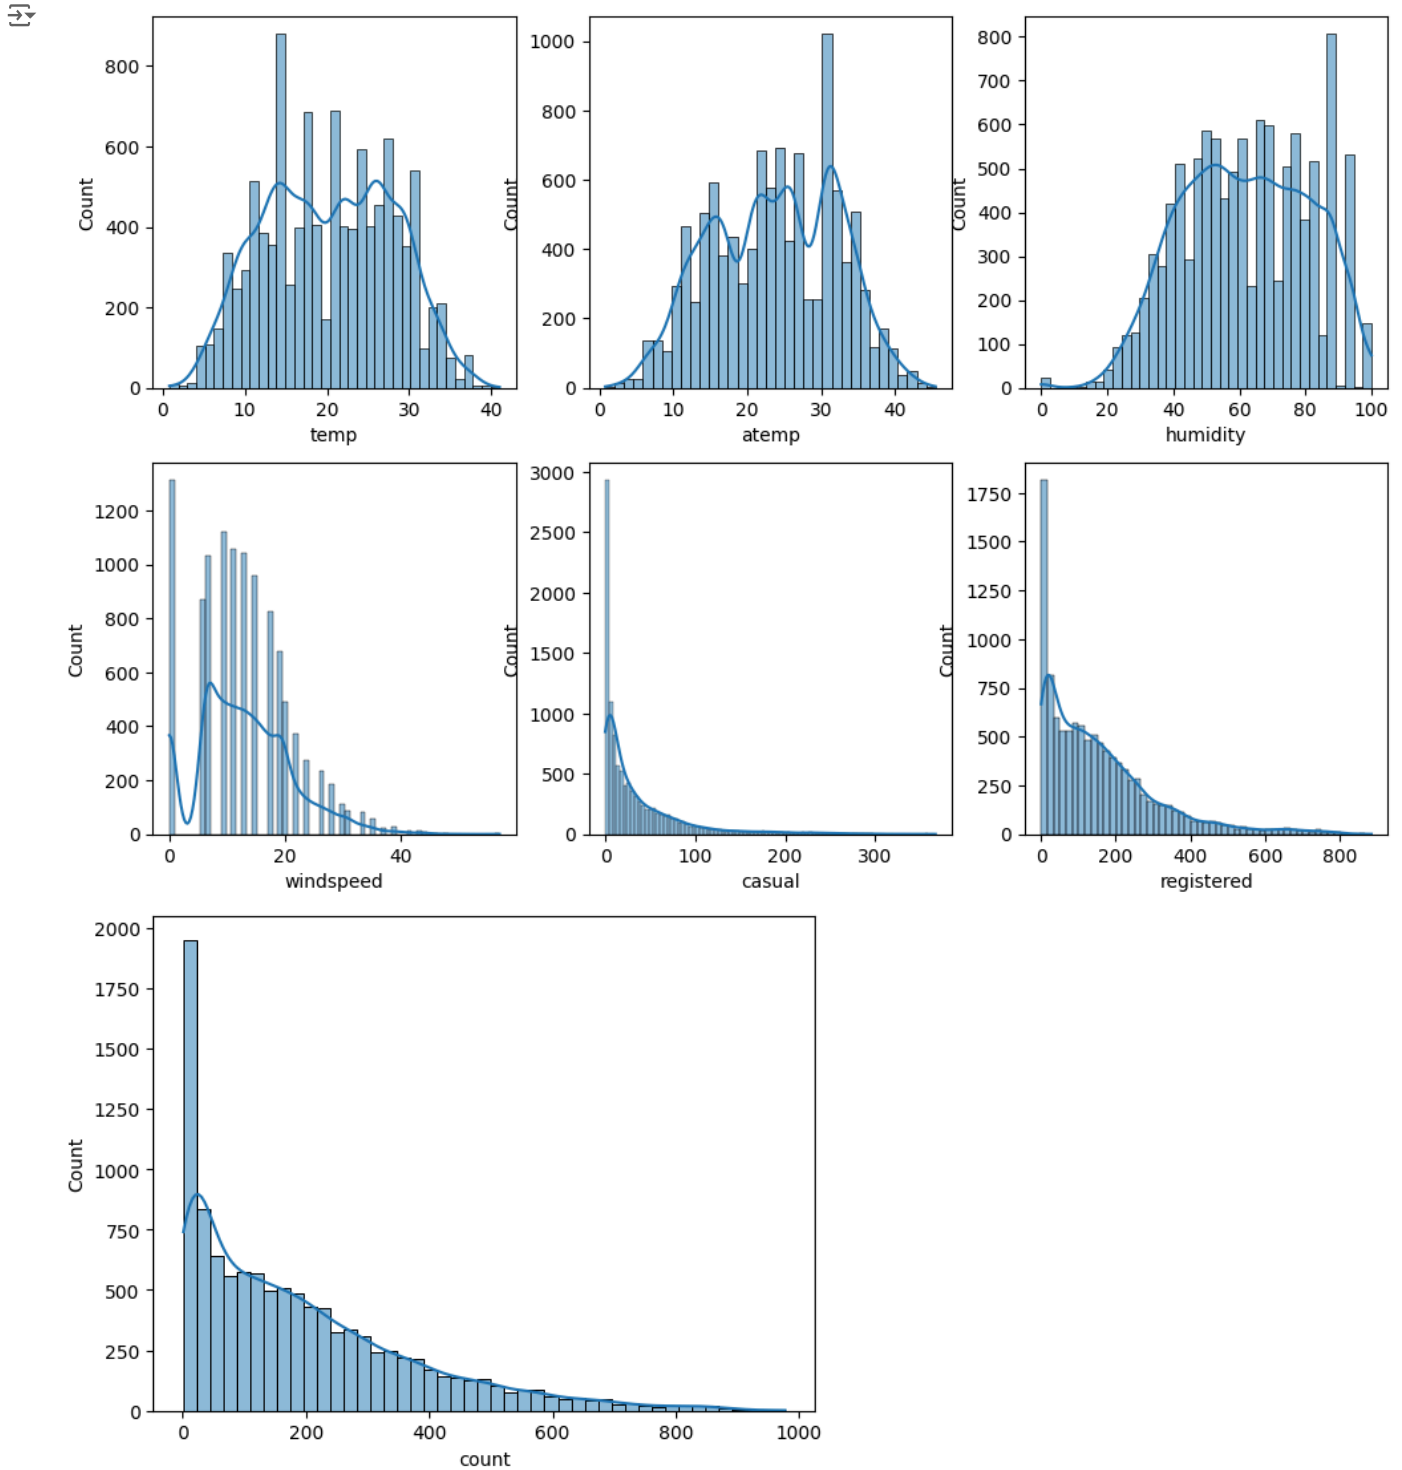
2011-01-01 00:00:00 2012-12-19 23:00:00		
	value	

variable	value	
holiday	0	10575
	1	311
season	1	2686
	2	2733
	3	2733
	4	2734
weather	1	7192
	2	2834
	3	859
	4	1
workingday	0	3474
	1	7412

```

1 # understanding the distribution for numerical variables
2 num_cols = ['temp', 'atemp', 'humidity', 'windspeed', 'casual', 'registered', 'count']
3
4 fig, axis = plt.subplots(nrows=2, ncols=3, figsize=(12, 8))
5
6 index = 0
7 for row in range(2):
8     for col in range(3):
9         sns.histplot(df[num_cols[index]], ax=axis[row, col], kde=True)
10        index += 1
11
12 plt.show()
13 sns.histplot(df[num_cols[-1]], kde=True)
14 plt.show()

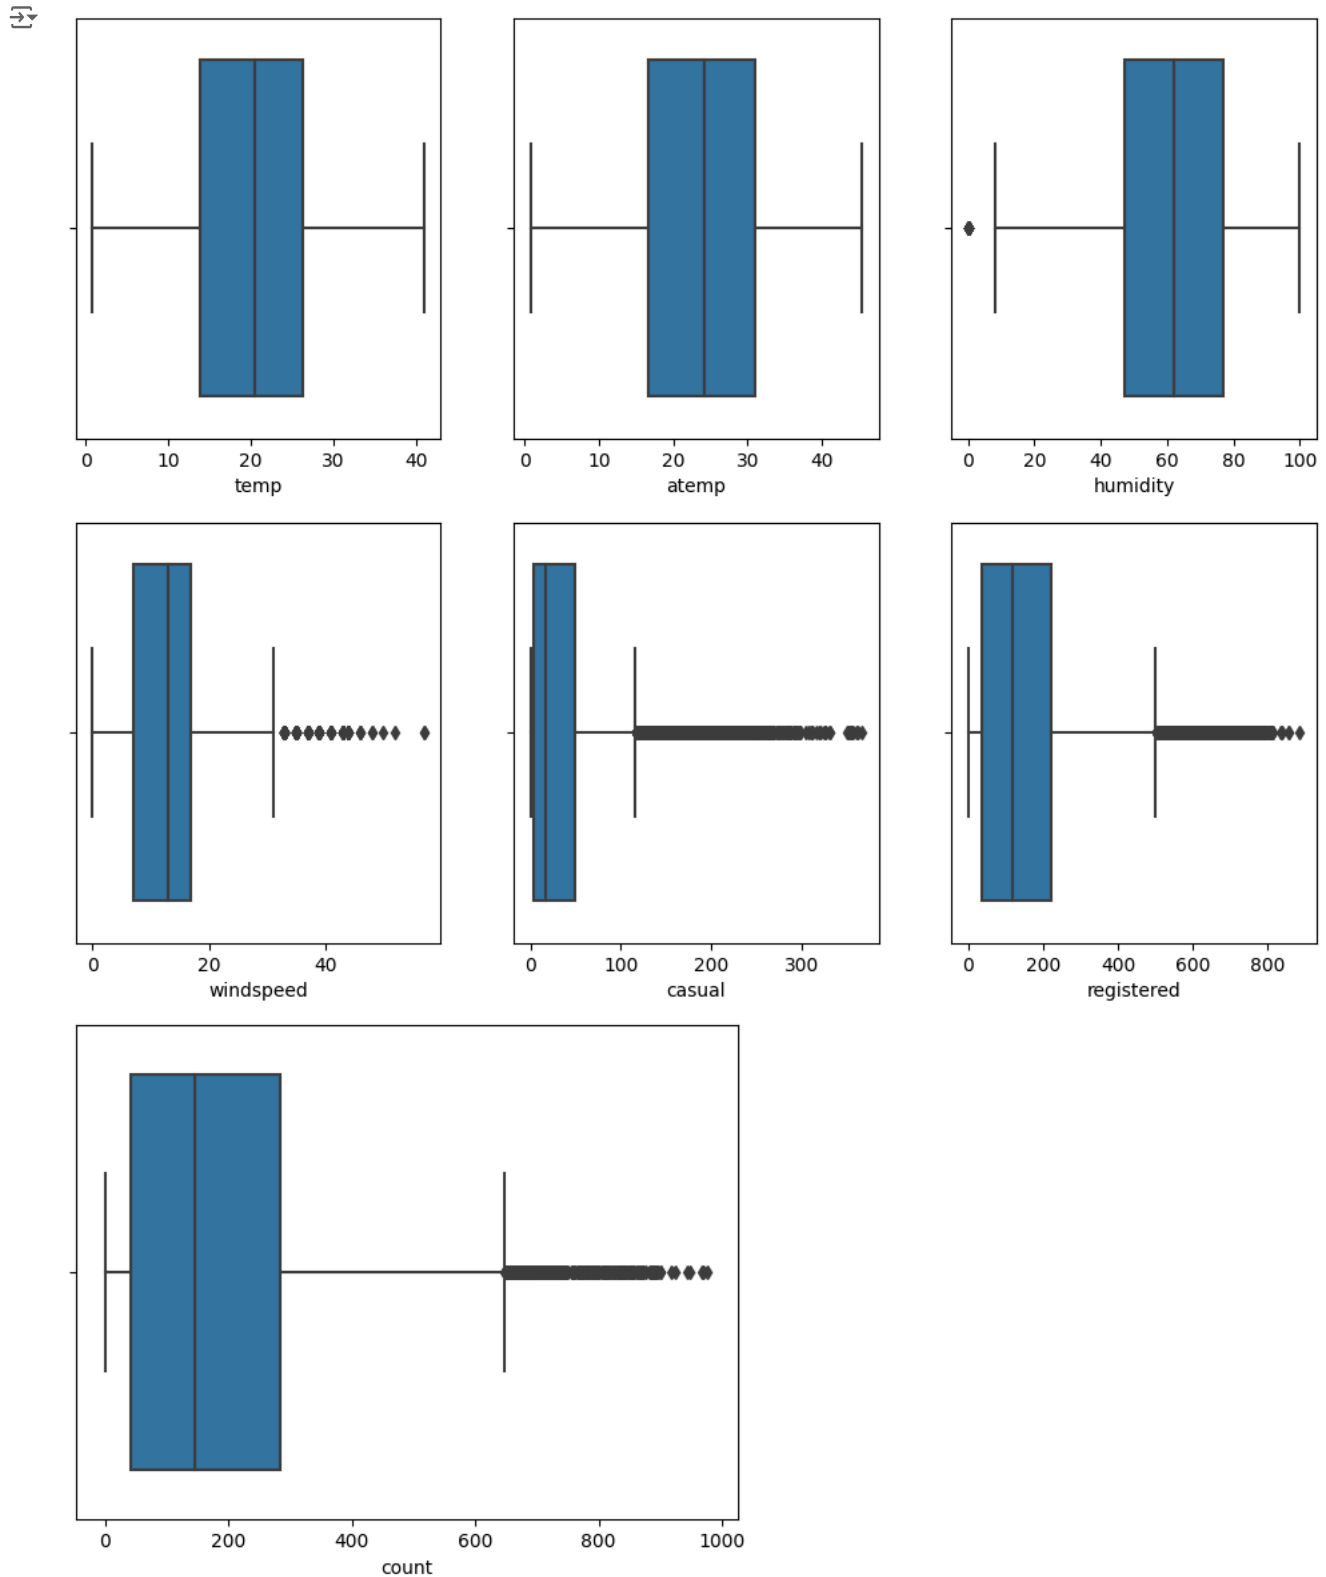
```



```

1 # plotting box plots to detect outliers in the data
2 fig, axis = plt.subplots(nrows=2, ncols=3, figsize=(12, 9))
3
4 index = 0
5 for row in range(2):
6     for col in range(3):
7         sns.boxplot(x=df[num_cols[index]], ax=axis[row, col])
8         index += 1
9
10 plt.show()
11 sns.boxplot(x=df[num_cols[-1]])
12 plt.show()


```



```

1 # countplot of each categorical column
2 fig, axis = plt.subplots(nrows=2, ncols=2, figsize=(8, 6))
3
4 index = 0
5 for row in range(2):
6     for col in range(2):
7         sns.countplot(data=df, x=cat_cols[index], ax=axis[row, col])
8         index += 1
9
10 plt.show()

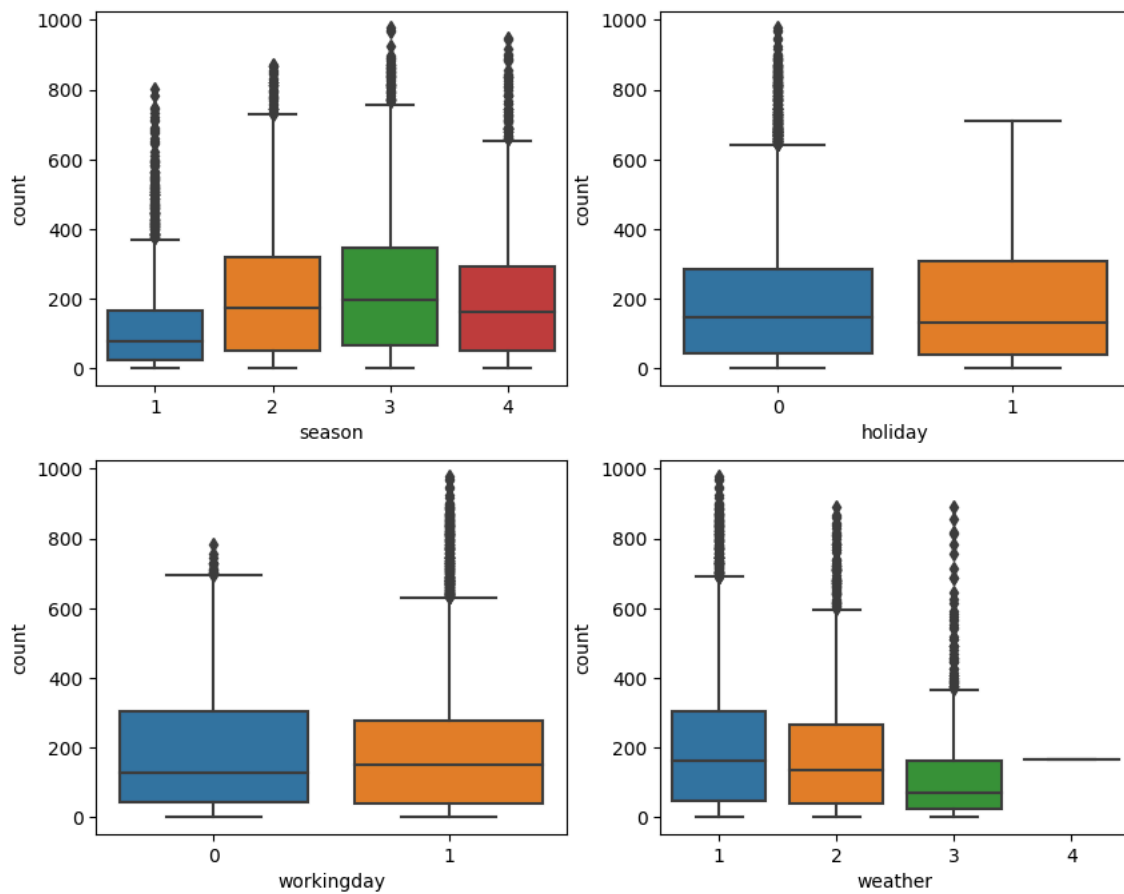
```



```

1 # plotting categorical variables against count using boxplots
2 fig, axis = plt.subplots(nrows=2, ncols=2, figsize=(10, 8))
3
4 index = 0
5 for row in range(2):
6     for col in range(2):
7         sns.boxplot(data=df, x=cat_cols[index], y='count', ax=axis[row, col])
8         index += 1
9
10 plt.show()

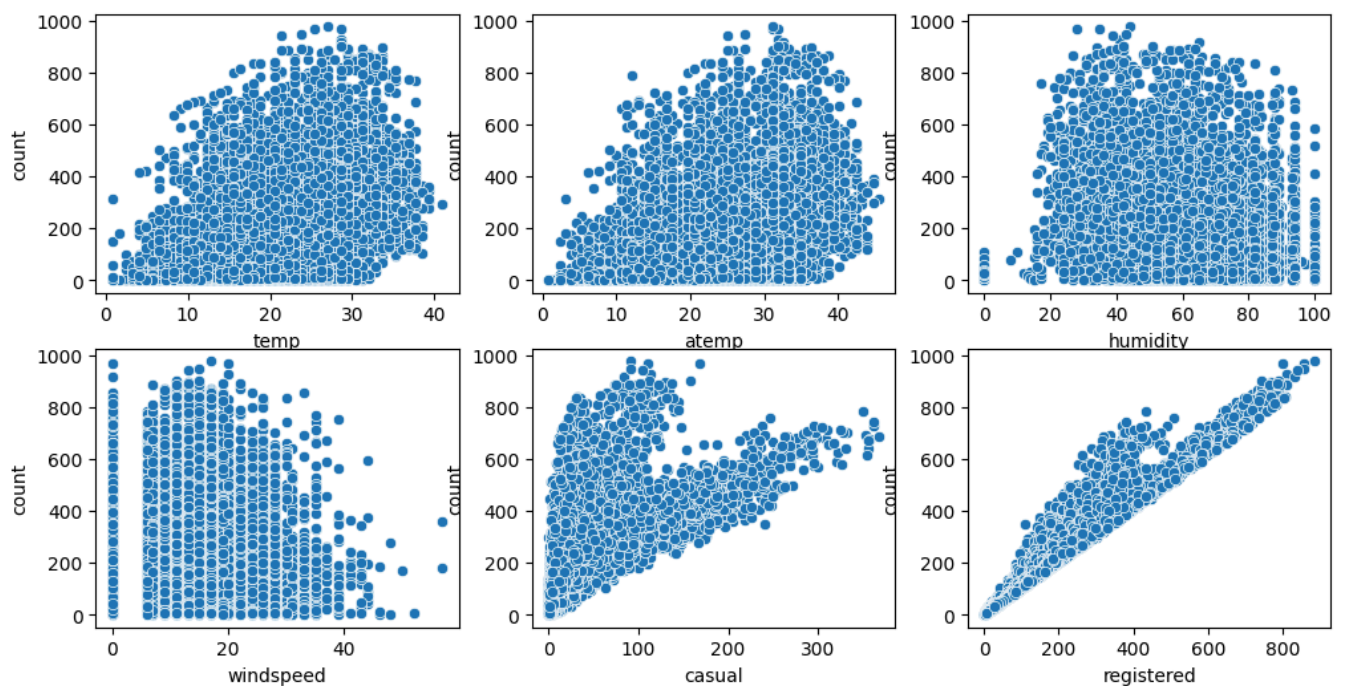
```



```

1 # plotting numerical variables against count using scatterplot
2 fig, axis = plt.subplots(nrows=2, ncols=3, figsize=(12, 6))
3
4 index = 0
5 for row in range(2):
6     for col in range(3):
7         sns.scatterplot(data=df, x=num_cols[index], y='count', ax=axis[row, col])
8         index += 1
9
10 plt.show()


```




```

1 # understanding the correlation between count and numerical variables
2 df.corr()['count']
3 sns.heatmap(df.corr(), annot=True)
4 plt.show()

```

 <ipython-input-22-b0729b22659f>:2: FutureWarning: The default value of numeric\_only in DataFrame.corr is deprecated. In df.corr()['count']

<ipython-input-22-b0729b22659f>:3: FutureWarning: The default value of numeric\_only in DataFrame.corr is deprecated. In sns.heatmap(df.corr(), annot=True)



```

1 data_table = pd.crosstab(df['season'], df['weather'])
2 print("Observed values:")
3 data_table

```

 Observed values:

weather	1	2	3	4
season				
1	1759	715	211	1
2	1801	708	224	0
3	1930	604	199	0
4	1702	807	225	0

```

1 val = stats.chi2_contingency(data_table)
2 print(val)
3 expected_values = val[3]
4 print(expected_values)
5 nrows, ncols = 4, 4
6 dof = (nrows-1)*(ncols-1)
7 print("degrees of freedom: ", dof)
8 alpha = 0.05
9
10
11 chi_sqr = sum([(o-e)**2/e for o, e in zip(data_table.values, expected_values)])
12 chi_sqr_statistic = chi_sqr[0] + chi_sqr[1]
13 print("chi-square test statistic: ", chi_sqr_statistic)
14
15 critical_val = stats.chi2.ppf(q=1-alpha, df=dof)
16 print(f"critical value: {critical_val}")
17
18 p_val = 1-stats.chi2.cdf(x=chi_sqr_statistic, df=dof)
19 print(f"p-value: {p_val}")
20
21 if p_val <= alpha:
22     print("\nSince p-value is less than the alpha 0.05, We reject the Null Hypothesis. Meaning that\
23     Weather is dependent on the season.")
24 else:
25     print("Since p-value is greater than the alpha 0.05, We do not reject the Null Hypothesis")

```

```

Chi2ContingencyResult(statistic=49.158655596893624, pvalue=1.549925073686492e-07, dof=9, expected_freq=array([[1.7745463
[1.80559765e+03, 7.11493845e+02, 2.15657450e+02, 2.51056403e-01],
[1.80559765e+03, 7.11493845e+02, 2.15657450e+02, 2.51056403e-01],
[1.80625831e+03, 7.11754180e+02, 2.15736359e+02, 2.51148264e-01]]))
[[1.77454639e+03 6.99258130e+02 2.11948742e+02 2.46738931e-01]
[1.80559765e+03 7.11493845e+02 2.15657450e+02 2.51056403e-01]
[1.80559765e+03 7.11493845e+02 2.15657450e+02 2.51056403e-01]
[1.80625831e+03 7.11754180e+02 2.15736359e+02 2.51148264e-01]]
degrees of freedom: 9
chi-square test statistic: 44.09441248632364
critical value: 16.918977604620448
p-value: 1.3560001579371317e-06

```

Since p-value is less than the alpha 0.05, We reject the Null Hypothesis. Meaning that Weather is dependent on the se

```

1 data_group1 = df[df['workingday']==0]['count'].values
2 data_group2 = df[df['workingday']==1]['count'].values
3
4 print(np.var(data_group1), np.var(data_group2))
5 np.var(data_group2)// np.var(data_group1)

```

```

30171.346098942427 34040.69710674686
1.0

```

```

1 stats.ttest_ind(a=data_group1, b=data_group2, equal_var=True)

```

```

Ttest_indResult(statistic=-1.2096277376026694, pvalue=0.22644804226361348)

```

```

1 # defining the data groups for the ANOVA
2 from statsmodels.graphics.gofplots import qqplot
3 gp1 = df[df['weather']==1]['count'].values
4 gp2 = df[df['weather']==2]['count'].values
5 gp3 = df[df['weather']==3]['count'].values
6 gp4 = df[df['weather']==4]['count'].values
7
8 gp5 = df[df['season']==1]['count'].values
9 gp6 = df[df['season']==2]['count'].values
10 gp7 = df[df['season']==3]['count'].values
11 gp8 = df[df['season']==4]['count'].values
12 groups=[gp1,gp2,gp3,gp4,gp5,gp6,gp7,gp8]
13
14


```

```

1 fig, axis = plt.subplots(nrows=4, ncols=2, figsize=(8, 8))
2
3 index = 0
4 for row in range(4):
5     for col in range(2):
6         sns.histplot(groups[index], ax=axis[row, col], kde=True)
7         index += 1
8
9 plt.show()

```





```

1
2 index = 0
3 for row in range(4):
4     for col in range(2):
5         qqplot(groups[index], line="s")
6         index += 1
7
8 plt.show()

```